Multi-Agent Reinforcement Learning for Traffic Signal Optimization

09.02.2026

—

Midhun Chandrasekhar
B Tech, MBA
Dubai
United Arab Emerites

1. Introduction

Traffic signal control has been a long-standing research problem within intelligent transportation systems. With increasing urbanization and traffic complexity, traditional control mechanisms have struggled to adapt to dynamic and stochastic traffic environments.

Recent advances in Artificial Intelligence, particularly **Reinforcement Learning (RL)** and **Multi-Agent Reinforcement Learning (MARL)**, have introduced promising approaches for adaptive and decentralized traffic management.

This section reviews existing work across three key paradigms:

- Traditional traffic control systems
- Reinforcement learning-based approaches
- Multi-agent and decentralized traffic optimization

2. Traditional Traffic Signal Control

2.1 Fixed-Time Control Systems

Fixed-time systems operate on predefined signal cycles based on historical traffic data.

Advantages:

- Simple to implement
- Low computational requirements

Limitations:

- No adaptability to real-time conditions
- Inefficient under fluctuating traffic patterns

2.2 Actuated and Adaptive Systems

Adaptive systems use sensors (loop detectors, cameras) to adjust signal timings dynamically.

Examples include:

- SCOOT (Split Cycle Offset Optimization Technique)

- SCATS (Sydney Coordinated Adaptive Traffic System)

Advantages:

- Respond to real-time traffic conditions
- Improved efficiency over fixed systems

Limitations:

- Centralized architecture
- Limited scalability
- Reactive rather than predictive behavior

3. Reinforcement Learning in Traffic Control

Reinforcement Learning (RL) models traffic control as a sequential decision-making problem, where an agent learns optimal policies through interaction with the environment.

3.1 Single-Agent RL Approaches

Early research focused on **single intersection control**, where one agent optimizes signal timing.

Common techniques:

- Q-Learning
- Deep Q Networks (DQN)

Key Findings:

- RL can outperform fixed-time systems
- Effective in isolated intersections

Limitations:

- Poor scalability to multi-intersection networks
- Lack of coordination between intersections

3.2 Deep Reinforcement Learning (DRL)

With the advancement of deep learning, DRL approaches have been used to handle high-dimensional traffic states.

Examples:

- Use of neural networks for state representation
- Learning complex traffic patterns

Advantages:

- Handles large state spaces
- Learns non-linear traffic dynamics

Challenges:

- High training complexity
- Stability issues
- Requires significant computational resources

4. Multi-Agent Reinforcement Learning (MARL)

To address scalability and coordination challenges, researchers have explored **Multi-Agent Reinforcement Learning**, where each intersection is modeled as an independent agent.

4.1 Decentralized Control

In MARL systems:

- Each intersection operates independently
- Agents learn based on local observations
- Global coordination emerges indirectly

Advantages:

- Scalable to large networks
- Reduced dependency on central controllers
- Better fault tolerance

4.2 Coordination Mechanisms

Research explores different coordination strategies:

- Independent learning (no explicit communication)
- Cooperative learning (shared rewards)

- Communication-based MARL

Observations:

- Independent learning is simpler but less optimal
- Cooperative approaches improve performance but increase complexity

4.3 Challenges in MARL

Despite its potential, MARL introduces several challenges:

- Non-stationary environment (agents constantly adapting)
- Credit assignment problem (which agent contributed to success?)
- Training instability
- Scalability vs coordination trade-off

5. Simulation Environments in Traffic Research

Simulation plays a critical role in evaluating traffic control systems.

Common tools:

- SUMO (Simulation of Urban Mobility)
- CityFlow
- Custom grid-based simulators

Importance:

- Enables controlled experimentation
- Allows comparison across models
- Facilitates reproducibility

However, many studies lack **standardized evaluation frameworks**, making comparisons difficult.

6. Identified Research Gaps

Based on the literature, several gaps remain:

6.1 Scalability Limitations

- Many RL approaches do not scale effectively beyond small networks

6.2 Centralization Bottlenecks

- Adaptive systems often rely on centralized architectures

6.3 Limited Real-World Alignment

- Simulations do not fully capture real traffic complexity

6.4 Coordination Complexity in MARL

- Effective decentralized coordination remains an open problem

7. Positioning of This Research

This research aims to address the above gaps through:

- A **decentralized MARL framework**
- A **controlled simulation environment** for consistent evaluation
- Comparative analysis between **baseline and adaptive systems**
- Exploration of **emergent coordination without explicit communication**

Unlike purely theoretical studies, this work emphasizes:

- Practical implementation
- Measurable performance improvements
- Scalability through decentralized design

8. Summary

The literature demonstrates that while reinforcement learning has significantly improved traffic signal optimization, challenges remain in scalability, coordination, and real-world applicability.

Multi-Agent Reinforcement Learning presents a promising direction, but requires further exploration to achieve stable, scalable, and efficient solutions.

This research builds upon these foundations to develop a **practical, decentralized, and adaptive traffic control system**, with a focus on real-world relevance and experimental validation.